

1

## Source Digitization and OCR-Derived Corpus

Physical 9-volume set (2nd ed., completed 2019–2020) → page photographs → PDF → MinerU extraction  
48 OCR-derived Markdown files · 759,186 lines · 16.6 M characters · rights-restricted source



2

## Corpus Profiling & Heterogeneity Characterization

181 HTML table/block remnants · 18,179 joined-entry signals · 24,239 malformed-heading signals  
Non-entry content: front matter, table-of-contents, indices, image markers · Noise quantified prior to extraction



3

## Entry-Level Structured Extraction & Provenance

97,527 parsed records · Per-record provenance: source file + character/line span  
Doc-type stratification (5 classes): clean | noisy | redirect | index-like | mixed/unknown

**Structural checks** | Record–span correspondence: 97,527 / 97,527 | Field consistency flags: 2

The two field flags identify merged-boundary cases in which a heading has no extracted value

Each structured record retains source-file and character/line-span coordinates



### Task A: Entry Segmentation

600-item human gold (2 annotators + adjudication)  
File-level train/dev/test split; held-out test evaluation  
Four baselines: heuristic, rule, char TF-IDF SVM,  
frozen BGE embeddings + logistic regression  
Bootstrap 95% CIs · exact McNemar tests · noise slices

### Task B: Field Robustness

Heading presence: 128 spans; exact values: 84,599 records  
Strict parser compared with a relaxed parser reference  
Fields: composition, administration, actions,  
indications (recall loss quantification)  
Per-noise-type error slice analysis