

OCR-Derived Corpus

1

Zhongyi Fangji Da Cidian, 2nd ed. (2019) · 48 Markdown files from 9 volumes
MinerU OCR/PDF extractor · 759,186 lines · 16.6 M characters · Rights-restricted source text



Corpus Profiling & Heterogeneity Characterization

2

181 HTML table/block remnants · 18,179 suspected glued-entry signals · 24,239 field-heading noise signals
Non-entry content: front matter, table-of-contents, indices, image markers · Noise quantified prior to extraction



Entry-Level Structured Extraction & Provenance

3

97,527 parsed records · Per-record provenance: source file + character/line span
Doc-type stratification (5 classes): clean | noisy | redirect | index-like | mixed/unknown

Quality Verification | Span fidelity: 97,527 / 97,527 pass (100%) | Field self-check: 2 flagged (both explainable)

Confirms provenance traceability of extraction; does not assert semantic correctness

Each record: raw_text field exactly matches the source character span declared in provenance



Task A: Entry Segmentation

300-item human gold (2 annotators + adjudication)
Cohen's kappa inter-annotator agreement
Silver consensus (full corpus; agree/disagree split)
3 baseline families: heuristic, rule-based,
char 3-gram naive Bayes · noise-slice evaluation

Task B: Field Robustness

Strict vs. relaxed field-heading parsing comparison
Recall degradation under OCR/Markdown artifacts
Fields: composition, administration, actions,
indications (recall loss quantification)
Per-noise-type error slice analysis