

1

Source Digitization and OCR-Derived Corpus

Physical 9-volume set (2nd ed., 2019) → page photographs → PDF → MinerU extraction
48 OCR-derived Markdown files · 759,186 lines · 16.6 M characters · rights-restricted source



2

Corpus Profiling & Heterogeneity Characterization

181 HTML table/block remnants · 18,179 suspected glued-entry signals · 24,239 field-heading noise signals
Non-entry content: front matter, table-of-contents, indices, image markers · Noise quantified prior to extraction



3

Entry-Level Structured Extraction & Provenance

97,527 parsed records · Per-record provenance: source file + character/line span
Doc-type stratification (5 classes): clean | noisy | redirect | index-like | mixed/unknown

Quality Verification | Span fidelity: 97,527 / 97,527 pass (100%) | Field self-check: 2 flagged (both explainable)

Confirms provenance traceability of extraction; does not assert semantic correctness

Each record: raw_text field exactly matches the source character span declared in provenance



Task A: Entry Segmentation

600-item human gold (2 annotators + adjudication)
File-level train/dev/test split; held-out test evaluation
Four baselines: heuristic, rule, char TF-IDF SVM,
frozen BGE embeddings + logistic regression
Bootstrap 95% CIs · exact McNemar tests · noise slices

Task B: Field Robustness

Strict vs. relaxed field-heading parsing comparison
Recall degradation under OCR/Markdown artifacts
Fields: composition, administration, actions,
indications (recall loss quantification)
Per-noise-type error slice analysis